

COE 347 HW 4 17 February 25

1. $u'' = f(x) = 4\pi^2 \cos(4\pi x)$

at $[0, 1]$

$$u''(x_i) h^2 = u_{i+1} - 2u_i + u_{i-1} + Ch^4 u^{(4)}(x_i) + h \cdot \text{O.T.} \quad (1)$$

$C = ?$

$$u^{(3)} = -16\pi^3 \sin(4\pi x)$$

$$u^{(4)} = -64\pi^4 \cos(4\pi x)$$

$$u_{i+1} = u(x_i + h) = u_i + h u'_i + \frac{h^2}{2} u''_i + \frac{h^3}{6} u'''_i + \dots + \frac{h^n}{n!} u^{(n)}_i$$

$$u_{i-1} = u(x_i - h) = u_i - h u'_i + \frac{h^2}{2} u''_i - \frac{h^3}{6} u'''_i + \dots - \frac{(-1)^n h^n}{n!} u^{(n)}_i$$

$$u_{i+1} + u_{i-1} = 2 \left[u_i + \frac{h^2}{2} u''_i + \frac{h^4}{4!} u'''_i - \frac{h^{2n} u^{(2n)}_i}{(2n)!} \right]$$

$$u_{i+1} + u_{i-1} - 2u_i = \underbrace{h^2 u''_i}_{\frac{h^4}{12} u'''_i} + \underbrace{\frac{h^4}{12} u'''_i}_{\frac{h^{2n}}{(2n)!} u^{(2n)}_i} + \dots + \underbrace{\frac{h^{2n} u^{(2n)}_i}{(2n)!}}_{\text{H.O.T.}} \quad (2)$$

$$\therefore u''_i h^2 = u_{i+1} - 2u_i + u_{i-1} - \left[\frac{h^4}{12} u'''_i + \dots \right] \text{H.O.T.} \quad (3)$$

→ compare (3) with (1) →

$$\boxed{C = -\frac{1}{12}}$$

p2_p3

Exported 2/17/2026, 11:05:19 AM

2

In [1]:

```

import numpy as np
import matplotlib.pyplot as plt

def solve_bvp(N=10):
    L = 1.0
    h = L / (N + 1)
    x = np.linspace(0, 1, N + 2)

    f = lambda x: (4*np.pi)**2 * np.cos(4*np.pi*x)

    u = np.zeros(N + 2)
    u[0] = 0
    u[-1] = 2

    A = np.zeros((N, N))
    b = np.zeros(N)

    for i in range(N):
        A[i, i] = -2.0
        if i > 0:
            A[i, i-1] = 1.0
        if i < N-1:
            A[i, i+1] = 1.0

        b[i] = h**2 * f(x[i+1])

    # Inject boundary conditions
    b[0] -= u[0]
    b[-1] -= u[-1]

    u[1:-1] = np.linalg.solve(A, b)
    return x, u

def exact_solution(x):
    """
    Exact solution to  $u''(x) = (4\pi)^2 \cos(4\pi x)$  with  $u(0)=0$ ,  $u(1)=2$ 

    The general solution is:  $u(x) = -\cos(4\pi x) + Ax + B$ 
    Applying boundary conditions:
     $u(0) = -1 + B = 0 \rightarrow B = 1$ 
     $u(1) = -\cos(4\pi) + A + 1 = -1 + A + 1 = A = 2 \rightarrow A = 2$ 

    So exact solution:  $u(x) = -\cos(4\pi x) + 2x + 1$ 
    """
    return -np.cos(4*np.pi*x) + 2*x + 1

```

The document was converted with Code-Format: <https://code-format.com/>

```
def main(N=10, analytic=True, print_results=True, plot=True):
    # Solve with N=10

    # N = 10
    x, u_numerical = solve_bvp(N)

    # Calculate exact solution
    u_exact = exact_solution(x)

    # Print results
    if print_results:
        print(f"Solution with N={N} interior points (h=
{1/(N+1)}:.4f):")
        print("-" * 60)
        print(f"\t{x[:10]} {'Numerical':>15} {'Exact
(Analytically)':>15} {'Error':>15}")
        print("-" * 60)
        for i in range(1, N+1): # i from 1 to 10
            error = abs(u_numerical[i] - u_exact[i])
            print(f"\t{i:3d} {x[i]:10.6f} {u_numerical[i]:18.12f}
{u_exact[i]:18.12f} {error:15.6e}")
        print("=" * 60)

try:
    # Load reference solution (N=10000)
    ref_data = np.loadtxt('solutionA_N10000.dat')
    if ref_data is not None:
        x_ref = ref_data[:, 0]
        u_ref = ref_data[:, 1]

        # Interpolate reference solution to our grid
        from scipy.interpolate import interp1d
        interp_func = interp1d(x_ref, u_ref, kind='linear')
        u_ref_interp = interp_func(x)

        # Calculate error compared to reference
        error_ref = np.max(np.abs(u_numerical - u_ref_interp))
        print(f"\nMaximum error compared to N=10000 reference:
{error_ref:.6e}")

        # Plot results
        if plot:
            plt.figure(figsize=(10, 6))
            assert x.shape == u_numerical.shape, "x and
u_numerical must have the same shape"
            plt.scatter(x, u_numerical, marker='o', color='b',
s=12, label=f'Numerical (N={N})')
```

```

        plt.scatter(x_ref, u_ref, marker='*', color='g', s=12,
alpha = .5, label='Exact/Reference (N=10000)', linewidth=2)
        if analytic:
            plt.scatter(x, u_exact, marker='x', color='r',
s=6, label='Analytic solution', linewidth=2)

            plt.xlabel('x', fontsize=12)
            plt.ylabel('u(x)', fontsize=12)
            plt.title('Solution of $u''(x) = (4\pi)^2
\cos(4\pi x)$ with $u(0)=0, u(1)=2$ (provided data)', fontsize=14)
            plt.legend(fontsize=12)

except:
    print("\nNote: Reference file 'solutionA_N10000.dat' not found
for comparison.")
    print("Using exact solution for validation instead.")

    for i in range(len(x)):
        error = abs(u_numerical[i] - u_exact[i])
        print(f"{x[i]:10.4f} {u_numerical[i]:15.6f}
{u_exact[i]:15.6f} {error:15.6f}")

    # Calculate and print maximum error
    max_error = np.max(np.abs(u_numerical - u_exact))
    print("\n" + "=" * 60)
    print(f"Maximum error: {max_error:.6e}")
    print("=" * 60)

    # Plot results
    plt.figure(figsize=(10, 6))
    plt.scatter(x, u_numerical, 'bo-', label=f'Numerical (N={N})',
markersize=6)
    plt.plot(x, u_exact, 'r-', label='Exact Analytic Solution',
linewidth=2)
    plt.xlabel('x', fontsize=12)
    plt.ylabel('u(x)', fontsize=12)
    plt.title('Solution of $u''(x) = (4\pi)^2 \cos(4\pi x)$
with $u(0)=0, u(1)=2$ (analytical solution)', fontsize=14)
    plt.legend(fontsize=12)
    plt.grid(True, alpha=0.3)
    plt.tight_layout()
    plt.show()

return x, u_numerical, u_ref, u_exact

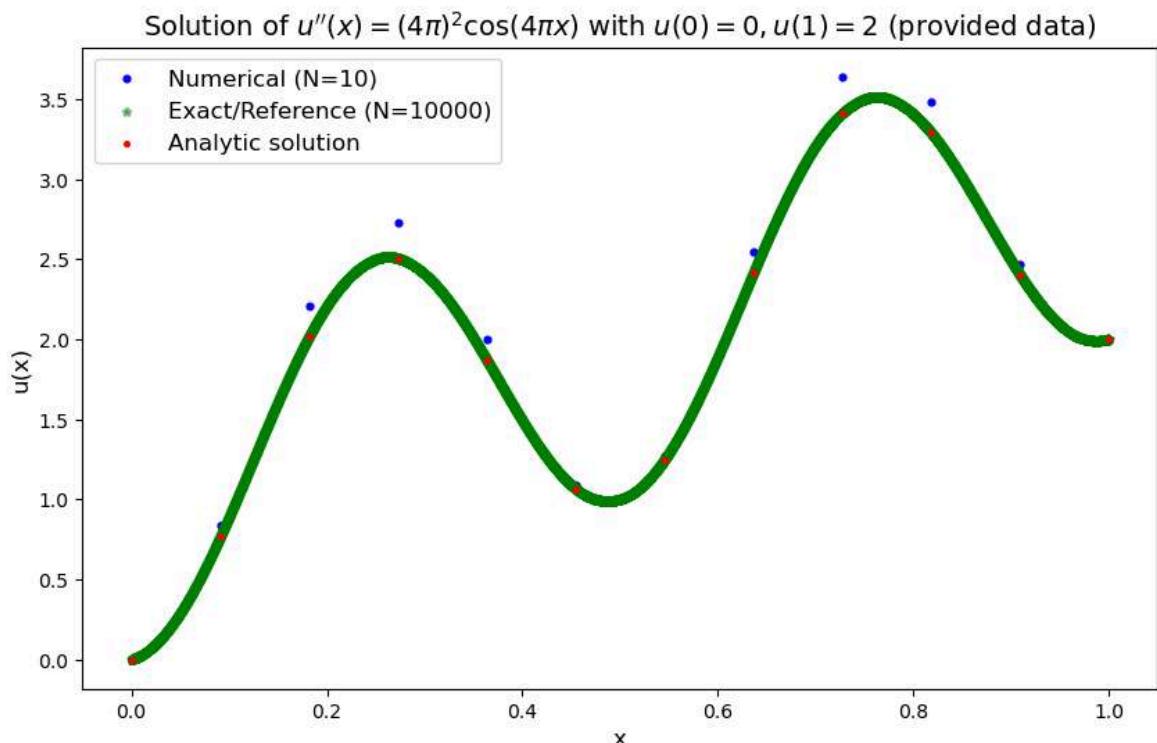
if __name__ == "__main__":
    main(N=10)

```

Solution with N=10 interior points ($h=0.0909$):

x	Numerical	Exact (Analytically)	Error
1 0.090909	0.834354009989	0.766403168816	6.795084e-02
2 0.181818	2.210854379067	2.018497097582	1.923573e-01
3 0.272727	2.732714565420	2.504947519069	2.277670e-01
4 0.363636	2.002367667450	1.869587565546	1.327801e-01
5 0.454545	1.086289707772	1.067837376260	1.845233e-02
6 0.545455	1.268107889591	1.249655558078	1.845233e-02
7 0.636364	2.547822212905	2.415042111001	1.327801e-01
8 0.727273	3.641805474511	3.414038428160	2.277670e-01
9 0.818182	3.483581651794	3.291224370309	1.923573e-01
10 0.909091	2.470717646353	2.402766805180	6.795084e-02

Maximum error compared to N=10000 reference: 2.277670e-01



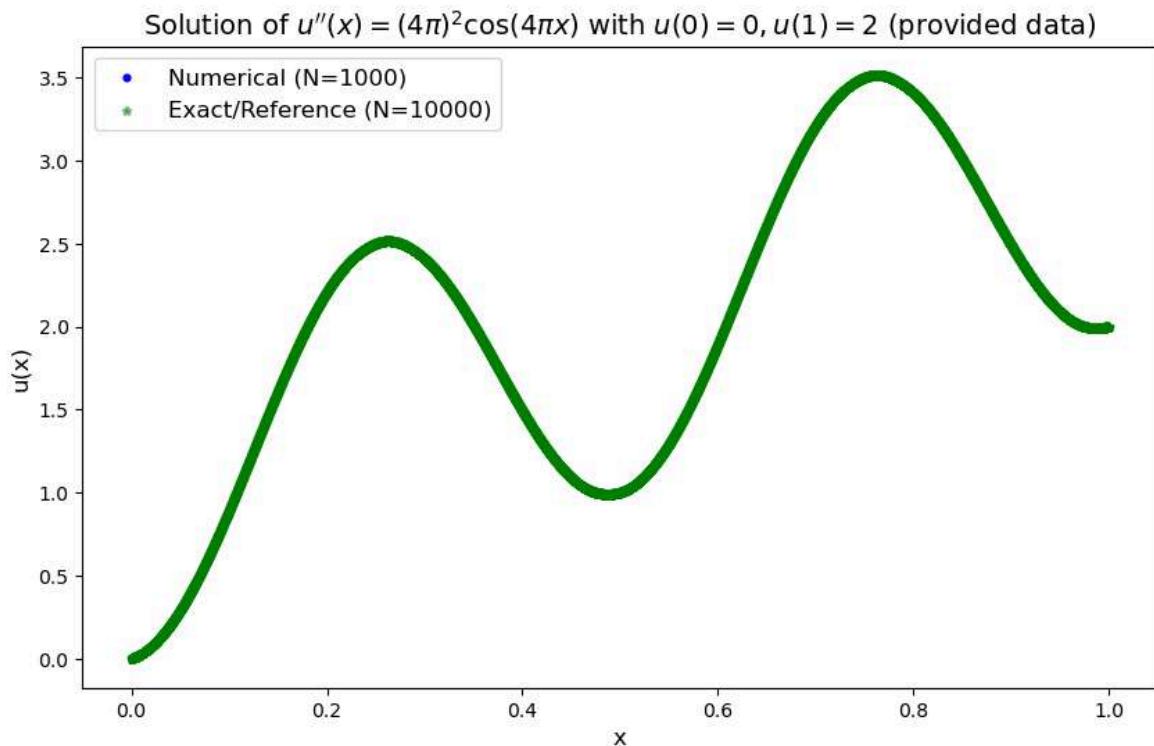
3

For confirmation that the numerical approximation converges to the true solution/the provided data, we can refine the mesh then compute the errors:

In []:

```
x, u_numerical, u_ref, u_exact = main(1000, analytic=False,  
print_results=False)
```

Maximum error compared to N=10000 reference: 2.615429e-05



In [3]:

```
from scipy.interpolate import CubicSpline
from scipy.optimize import curve_fit
import warnings; warnings.filterwarnings('ignore')

L = 1.
E = lambda u_num, u_ref: np.linalg.norm(u_num - u_ref)
e = lambda u_num, u_ref: np.linalg.norm(u_num - u_ref) / len(u_num)
N_values = [5, 10, 20, 40, 80, 160, 320, 640, 1280]
results = {
    'N': [],
    'h': [],
    'E': [],
    'e': [],
}

print(f"{'N':>6} {'h':>12} {'E (Frobenius)':>18} {'e (scaled)':>18}")

for N in N_values:
    x, u_numerical, u_ref, u_exact = main(N, analytic=False,
print_results=False, plot=False)
    x_ref = np.linspace(0, 1, 10000 + 2)
    uinterior_ref = CubicSpline(x_ref, u_ref)(x[1:-1])

    Error, error, = E(u_numerical[1:-1], uinterior_ref),
e(u_numerical[1:-1], uinterior_ref)
    h = L / (N + 1)

    results['N'].append(N)
    results['h'].append(h)
    results['E'].append(Error)
    results['e'].append(error)
    print(f"N:{6d} {h:12.6e}"
          f"{{Error:18.6e} {error:18.6e}}")

print("=" * 80)

for key in results:
    results[key] = np.array(results[key])

log_h = np.log(results['h'])
log_E = np.log(results['E'])
log_e = np.log(results['e'])
```

For E (Frobenius norm)

```

coeffs_E = np.polyfit(log_h, log_E, 1)
a_E = coeffs_E[0]
C_E = np.exp(coeffs_E[1])

# For e (scaled Frobenius norm)
coeffs_e = np.polyfit(log_h, log_e, 1)
a_e = coeffs_e[0]
C_e = np.exp(coeffs_e[1])

print(f"E (Frobenius) fit: E = {C_E:.6f} * h^{a_E:.4f}")
print(f"e (scaled) fit:     e = {C_e:.6f} * h^{a_e:.4f}")

# Create log-log plot
fig, ax = plt.subplots(figsize=(14, 10))

ax.loglog(1/results['h'], results['E'], 'bo', linewidth=2,
          markersize=8, label='Data E')
ax.loglog(1/results['h'], results['e'], 'bo', linewidth=2,
          markersize=8, label='Data e')

# Plot fitted curve
h_fit = np.logspace(np.log10(min(results['h'])),
                    np.log10(max(results['h'])), 100)
E_fit = C_E * (h_fit**a_E)
ax.loglog(1/h_fit, E_fit, 'r--', linewidth=2, label=f'E with Fit: a={a_E:.3f}')

e_fit = C_e * (h_fit**a_e)
ax.loglog(1/h_fit, e_fit, 'r--', linewidth=2, label=f'e with Fit: a={a_e:.3f}')

ax.set_xlabel('1/h', fontsize=12)
ax.set_ylabel('E (Frobenius norm)', fontsize=12)
ax.set_title('Frobenius Error vs 1/h', fontsize=14)
ax.grid(True, alpha=0.3)
ax.legend(fontsize=12)

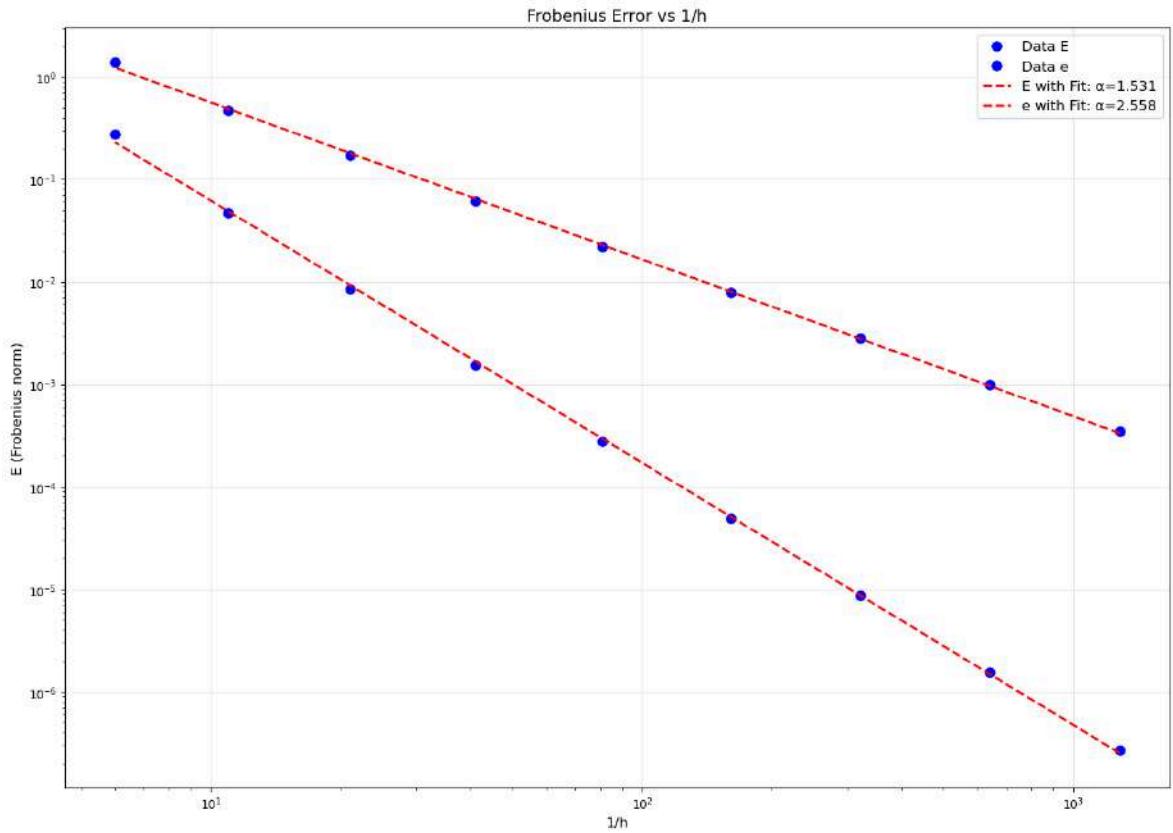
plt.tight_layout()
plt.savefig('error_convergence.png', dpi=150, bbox_inches='tight')
plt.show()

# Print convergence rates
print("\n" + "=" * 80)
print("CONVERGENCE RATE ANALYSIS:")
print("=" * 80)
print(f"Theoretical expectation: 2nd-order scheme → a ≈ 2")

```

```
print(f"Frobenius norm (E): α = {α_E:.4f} (expectation: 2.0)")  
print(f"Scaled norm (e):      α = {α_e:.4f} (expectation: 2.0)")
```

N	h	E (Frobenius)	e (scaled)
Maximum error compared to N=10000 reference: 6.932453e-01			
5	1.666667e-01	1.386490e+00	2.772981e-01
Maximum error compared to N=10000 reference: 2.277670e-01			
10	9.090909e-02	4.721594e-01	4.721594e-02
Maximum error compared to N=10000 reference: 6.042456e-02			
20	4.761905e-02	1.705179e-01	8.525896e-03
Maximum error compared to N=10000 reference: 1.570735e-02			
40	2.439024e-02	6.167992e-02	1.541998e-03
Maximum error compared to N=10000 reference: 4.014668e-03			
80	1.234568e-02	2.213354e-02	2.766693e-04
Maximum error compared to N=10000 reference: 1.015461e-03			
160	6.211180e-03	7.889783e-03	4.931114e-05
Maximum error compared to N=10000 reference: 2.553685e-04			
320	3.115265e-03	2.799704e-03	8.749076e-06
Maximum error compared to N=10000 reference: 6.397205e-05			
640	1.560062e-03	9.890482e-04	1.545388e-06
Maximum error compared to N=10000 reference: 1.597234e-05			
1280	7.806401e-04	3.457620e-04	2.701266e-07
=====			
=====			
E (Frobenius) fit: E = 19.084012 * h^1.5306			
e (scaled) fit: e = 22.446809 * h^2.5578			



CONVERGENCE RATE ANALYSIS:

Theoretical expectation: 2nd-order scheme $\rightarrow \alpha \approx 2$

Frobenius norm (E): $\alpha = 1.5306$ (expectation: 2.0)

Scaled norm (e): $\alpha = 2.5578$ (expectation: 2.0)

Approximately, E converges at order 1.5 and e converges at order 2.5. Theoretically, E should converge at 1.5, while e is expected to be of order 1 greater than the order of E: e is proportional to $1/(\text{number of interior points on mesh})$ which is proportional to the step size h.

For proof/explanation on expected vs theoretical orders of convergence:

Assume the pointwise discretization error satisfies

$$|u_h(x_i) - u(x_i)| = O(h^2).$$

Then the discrete ℓ^2 error is

$$E = \left(\sum_{i=1}^N O(h^4) \right)^{1/2} = O!(\sqrt{N}, h^2) = O(h^{3/2}),$$

since

$$N \sim \frac{1}{h}.$$

Next, define the scaled error

$$e = \frac{1}{N} \|u_h - u\|_{\ell^2}.$$

Because

$$\frac{1}{N} \sim h,$$

we obtain

$$e = h, E.$$

Therefore,

$$e = O!(h^{3/2+1}) = O!(h^{5/2}).$$

Hence, the observed convergence rate of order 2.5 for e is as theoretically expected.

$$4. u'(x_i)h = au_i + bu_{i+1} + cu_{i+2} + O(h^3) \quad (1)$$

$$u_{i+1} = u_i + hu'_i + \frac{h^2}{2}u''_i + \frac{h^3}{6}u^{(3)}_i + O(h^4) \quad (2)$$

$$\begin{aligned} u_{i+2} &= u_i + 2hu'_i + \frac{(2h)^2}{2}u''_i + \frac{(2h)^3}{6}u^{(3)}_i + O(h^4) \\ &= u_i + 2hu'_i + 2h^2u''_i + \frac{4}{3}h^3u^{(3)}_i + O(h^4) \end{aligned} \quad (3)$$

In equation 1, there are no second order terms (with u''_i), so, some combination of (2), (3), and u_i must eliminate 2nd order terms. u_i has no u''_i . So, some combination of (2), (3) eliminates u''_i :

$$\begin{aligned} (4u_{i+1} - u_{i+2}) &= 4u_i + 4hu'_i + 2h^2u''_i + \frac{2}{3}h^3u^{(3)}_i + O(h^4) \\ &\quad - u_i - 2hu'_i - 2h^2u''_i + \frac{4}{3}h^3u^{(3)}_i + O(h^4) \\ &= 3u_i + 2hu'_i + 2h^3u^{(3)}_i + O(h^4) \end{aligned}$$

So, we have:

$$\begin{array}{c} b=4d \\ u'(x_i)h = d(4u_{i+1} - u_{i+2}) + au_i + O(h^3) \end{array}$$

$$\begin{aligned} &= d(3u_i + 2hu'_i + 2h^3u^{(3)}_i) + au_i + O(h^3) \\ &= 3du_i + 2dh u'_i + 2dh^3u^{(3)}_i + O(h^3) \\ &\quad + au_i \end{aligned}$$

Since we have (1) $u'_i h$ on the LHS, $2d = 1$
and $d = 1/2$.

We also have (0) u_i on the LHS, so $3d + a = 0$.

$$\Rightarrow a = -3d = -3/2$$

Therefore: $b = 4d = 2$, $c = -d = -1/2$

3

$$u'(x_i) h = \underbrace{-\frac{3}{2} u_i}_a + \underbrace{2 u_{i+1}}_b - \underbrace{\frac{1}{2} u_{i+2}}_c + O(h^3) \quad (4)$$

$$\boxed{\begin{aligned} a &= -3/2 \\ b &= 2 \\ c &= -1/2 \end{aligned}}$$

To recheck:

RHS of eq (4):

$$\begin{aligned} &-\frac{3}{2} u_i + 2 \left[u_i + h u'_i + \frac{h^2}{2} u''_i + \frac{h^3}{3!} u^{(3)}_i + O(h^4) \right] \\ &- \frac{1}{2} \left[u_i + 2h u'_i + \frac{4h^2}{2} u''_i + \frac{8h^3}{3!} u^{(3)}_i + O(h^4) \right] \\ &= \left(-\frac{3}{2} + 2 - \frac{1}{2} \right) u_i \\ &\quad + (2h - h) u'_i \\ &\quad + \left(2h^2/2 - 4h^2/4 \right) u''_i \\ &\quad + O(h^3) \\ &= (0) u_i + h u'_i + (0) h^2 u''_i + O(h^3) = h u'_i = u'(x_i) h = \text{RHS. } \checkmark \end{aligned}$$

p5_p6

Exported 2/17/2026, 11:05:39 AM

5

In [7]:

```
import numpy as np
import matplotlib.pyplot as plt
import numpy as np

def solve_bvp(N=10, fd_order=2):
    L = 1.0
    h = L / (N + 1)
    x = np.linspace(0, 1, N + 2)

    f = lambda x: (4*np.pi)**2 * np.cos(4*np.pi*x)

    u = np.zeros(N + 2)

    u[-1] = 2

    A = np.zeros((N+1, N+1))
    b = np.zeros(N+1)

    for i in range(N+1):
        A[i, i] = -2.0
        if i > 0:
            A[i, i-1] = 1.0
        if i < N:
            A[i, i+1] = 1.0

        b[i] = h**2 * f(x[i])
    if fd_order == 2:
        b0 = 20*h
        A[0,0] = -3
        A[0,1] = 4
        A[0,2]=-1
    elif fd_order == 1:
        b0 = 10*h
        A[0,0] = -1
        A[0,1] = 1

    # Inject boundary conditions
    b[0] = b0
    b[-1] -= u[-1]

    u[0:-1] = np.linalg.solve(A, b)

    return x, u
```

```
def exact_solution(x):
```

```

"""
Exact solution to  $u''(x) = (4\pi)^2 \cos(4\pi x)$  with  $u'(0)=1$ ,  $u(1)=2$ 
 $u' = 4\pi \sin(4\pi x) + A$ 
The general solution is:  $u(x) = -\cos(4\pi x) + Ax + B$ 
Applying boundary conditions:
 $u'(0) = 0 + A = 10 \rightarrow A = 10$ 
 $u(1) = -\cos(4\pi) + 10 + B = -1 + 10 + B = 9 + B = 2 \rightarrow B = -7$ 

So exact solution:  $u(x) = -\cos(4\pi x) + 10x - 7$ 
"""

return -np.cos(4*np.pi*x) + 10*x - 7

```

```

def main(N=10, analytic=True, print_results=True, plot=True, size =
None, fd_order=2):
    # Solve with N=10

    # N = 10
    x, u_numerical = solve_bvp(N, fd_order = fd_order)

    # Calculate exact solution
    u_exact = exact_solution(x)

    # Print results
    if print_results:
        print(f"Solution with N={N} interior points (h=
{1/(N+1):.4f}):")
        print("-" * 60)
        print(f"{x[:10]} {'Numerical':>15} {'Exact
(Analytically)':>15} {'Error':>15}")
        print("-" * 60)
        for i in range(1, N+1):  # i from 1 to 10
            error = abs(u_numerical[i] - u_exact[i])
            print(f"{i:3d} {x[i]:10.6f} {u_numerical[i]:18.12f}
{u_exact[i]:18.12f} {error:15.6e}")
        print("=" * 60)

try:
    # Load reference solution (N=10000)
    ref_data = np.loadtxt('solutionB_N10000.dat')
    if ref_data is not None:
        x_ref = ref_data[:, 0]
        u_ref = ref_data[:, 1]

        # Interpolate reference solution to our grid
        from scipy.interpolate import interp1d
        interp_func = interp1d(x_ref, u_ref, kind='linear')

```

The document was converted with code-format.com/

```

# Calculate error compared to reference
error_ref = np.max(np.abs(u_numerical - u_ref_interp))
print(f"\nMaximum error compared to N=10000 reference:
{error_ref:.6e}")

# Plot results
if plot:
    plt.figure(figsize=(10, 6))
    assert x.shape == u_numerical.shape, "x and
u_numerical must have the same shape"
    if size:
        ms = size
    else:
        ms = 12
    plt.scatter(x, u_numerical, marker='o', color='b',
s=ms, label=f'Numerical (N={N})')

    plt.scatter(x_ref, u_ref, marker='*', color='g', s=12,
alpha = .5, label='Exact/Reference (N=10000)', linewidth=2)
    if analytic:
        plt.scatter(x, u_exact, marker='x', color='r',
s=6, label='Analytic solution', linewidth=2)

    plt.xlabel('x', fontsize=12)
    plt.ylabel('u(x)', fontsize=12)
    plt.title('Solution of $u''(x) = (4\pi)^2
\cos(4\pi x)$ with $u'(0)=10, u(1)=2$ (provided data)',
fontsize=14)
    plt.legend(fontsize=12)

except:
    print("\nNote: Reference file 'solutionA_N10000.dat' not found
for comparison.")
    print("Using exact solution for validation instead.")

for i in range(len(x)):
    error = abs(u_numerical[i] - u_exact[i])
    print(f"{x[i]:10.4f} {u_numerical[i]:15.6f}
{u_exact[i]:15.6f} {error:15.6f}")

# Calculate and print maximum error
max_error = np.max(np.abs(u_numerical - u_exact))
print("\n" + "=" * 60)
print(f"Maximum error: {max_error:.6e}")
print("=" * 60)

# Plot results
plt.figure(figsize=(10,6))

```

The original code was generated with Code-Format at <https://code-format.com/>

```

plt.scatter(x, u_numerical, 'bo-', label=f'Numerical (N={N})',
            markersize=6)
    plt.plot(x, u_exact, 'r-', label='Exact Analytic Solution',
              linewidth=2)
        plt.xlabel('x', fontsize=12)
        plt.ylabel('u(x)', fontsize=12)
        plt.title('Solution of $u''(x) = (4\pi)^2 \cos(4\pi x)$
with $u'(0)=10, u(1)=2$ (analytical solution)', fontsize=14)
        plt.legend(fontsize=12)
        plt.grid(True, alpha=0.3)
        plt.tight_layout()
        plt.show()
    return x, u_numerical, u_ref, u_exact

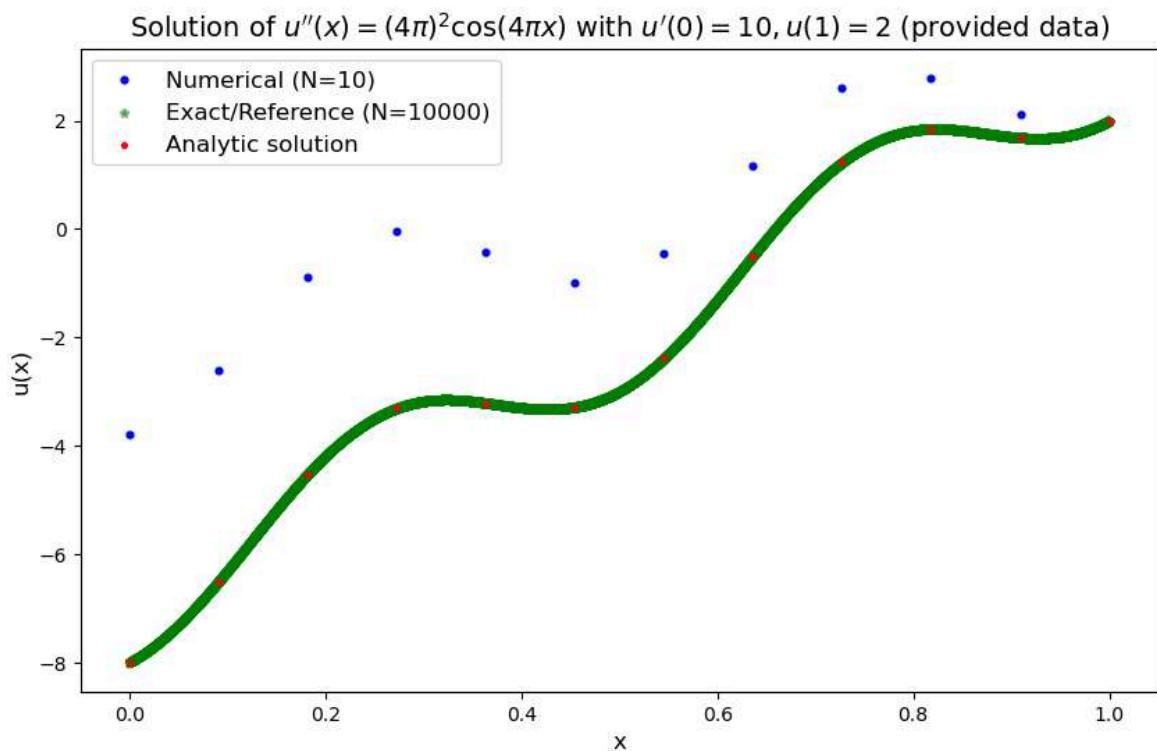
if __name__ == "__main__":
    main(N=10)

```

Solution with N=10 interior points (h=0.0909):

	x	Numerical	Exact (Analytically)	Error
1	0.090909	-2.623746776465	-6.506324103911	3.882577e+00
2	0.181818	-0.901436328743	-4.526957447873	3.625521e+00
3	0.272727	-0.033766063743	-3.313234299113	3.279468e+00
4	0.363636	-0.418302883068	-3.221321525363	2.803019e+00
5	0.454545	-0.988570764100	-3.295798987377	2.307228e+00
6	0.545455	-0.460942503637	-2.386708078286	1.925766e+00
7	0.636364	1.164581898323	-0.494048798090	1.658631e+00
8	0.727273	2.604375238575	1.232220246342	1.372155e+00
9	0.818182	2.791961494503	1.836678915763	9.552826e-01
10	0.909091	2.124907567708	1.675494077907	4.494135e-01

Maximum error compared to N=10000 reference: 4.196089e+00

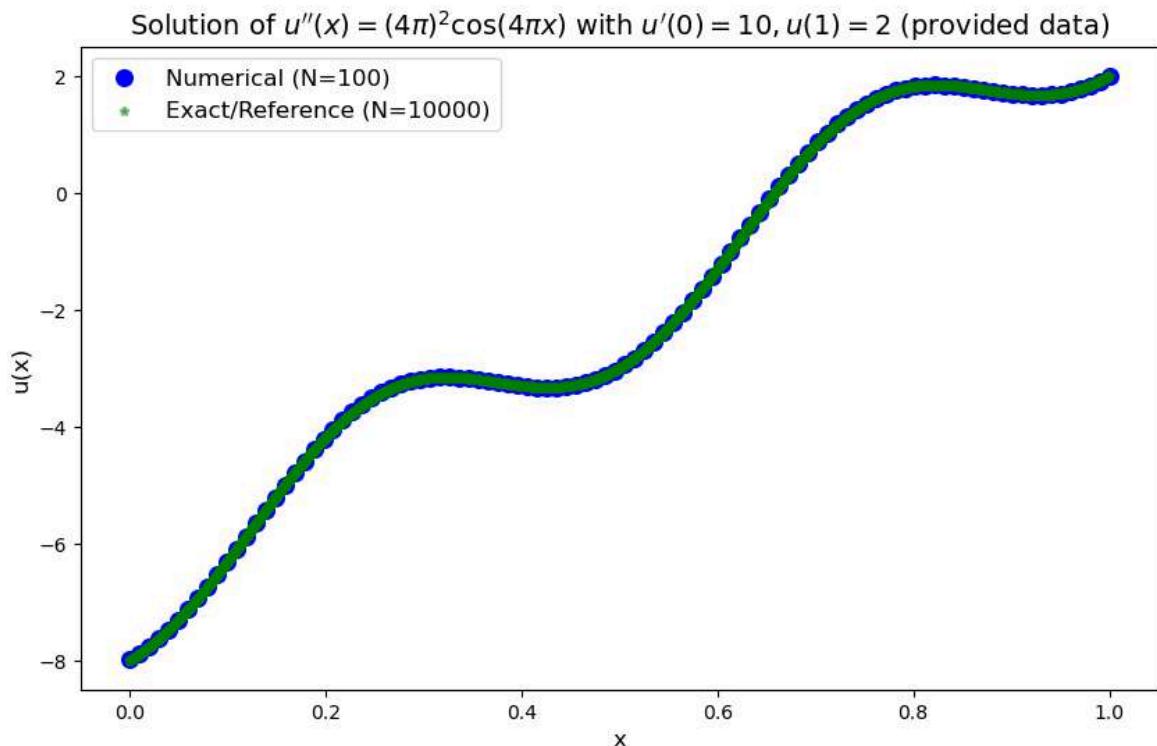


For confirmation that the numerical approximation converges to the true solution/the provided data, we can refine the mesh then compute the errors:

In [8]:

```
x, u_numerical, u_ref, u_exact = main(100, analytic=False,  
print_results=False, size = 70)
```

Maximum error compared to N=10000 reference: 7.204568e-03



6

In [13]:

```
from scipy.interpolate import CubicSpline
from scipy.optimize import curve_fit
import warnings; warnings.filterwarnings('ignore')

def fit_and_plot_convergence(fd_order=2, N_values = [5, 10, 20, 40,
80, 160, 320, 640, 1280]):
    L = 1.
    E = lambda u_num, u_ref: np.linalg.norm(u_num - u_ref)
    e = lambda u_num, u_ref: np.linalg.norm(u_num - u_ref) /
len(u_num)

    results = {
        'N': [],
        'h': [],
        'E': [],
        'e': [],
    }

    print(f"{'N':>6} {'h':>12} {'E (Frobenius)':>18} {'e
(scaled)':>18}")

    for N in N_values:
        x, u_numerical, u_ref, u_exact = main(N, analytic=False,
print_results=False, plot=False, fd_order=fd_order)
        x_ref = np.linspace(0, 1, 10000 + 2)
        uinterior_ref = CubicSpline(x_ref, u_ref)(x[1:-1])

        Error, error, = E(u_numerical[1:-1], uinterior_ref),
e(u_numerical[1:-1], uinterior_ref)
        h = L / (N + 1)

        results['N'].append(N)
        results['h'].append(h)
        results['E'].append(Error)
        results['e'].append(error)
        print(f"{{N:6d} {h:12.6e}""
f"{{Error:18.6e} {error:18.6e}}")

    print("=" * 80)

    for key in results:
        results[key] = np.array(results[key])

    log_h = np.log(results['h'])
    log_E = np.log(results['E'])
    log_e = np.log(results['e'])
```

The document was converted with Code-Format: <https://code-format.com/>

```
# For E (Frobenius norm)
coeffs_E = np.polyfit(log_h, log_E, 1)
a_E = coeffs_E[0]
C_E = np.exp(coeffs_E[1])

# For e (scaled Frobenius norm)
coeffs_e = np.polyfit(log_h, log_e, 1)
a_e = coeffs_e[0]
C_e = np.exp(coeffs_e[1])

print(f"E (Frobenius) fit: E = {C_E:.6f} * h^{a_E:.4f}")
print(f"e (scaled) fit:     e = {C_e:.6f} * h^{a_e:.4f}")

# Create log-log plot
fig, ax = plt.subplots(figsize=(14, 10))

    ax.loglog(1/results['h'], results['E'], 'bo', linewidth=2,
    markersize=8, label='Data E')
    ax.loglog(1/results['h'], results['e'], 'bo', linewidth=2,
    markersize=8, label='Data e')

# Plot fitted curve
h_fit = np.logspace(np.log10(min(results['h'])),
np.log10(max(results['h'])), 100)
E_fit = C_E * (h_fit**a_E)
ax.loglog(1/h_fit, E_fit, 'r--', linewidth=2, label=f'E with Fit:
a={a_E:.3f}')

e_fit = C_e * (h_fit**a_e)
ax.loglog(1/h_fit, e_fit, 'r--', linewidth=2, label=f'e with Fit:
a={a_e:.3f}')

ax.set_xlabel('1/h', fontsize=12)
ax.set_ylabel('E (Frobenius norm)', fontsize=12)
ax.set_title('Frobenius Error vs 1/h', fontsize=14)
ax.grid(True, alpha=0.3)
ax.legend(fontsize=12)

plt.tight_layout()
plt.savefig('error_convergence.png', dpi=150, bbox_inches='tight')
plt.show()

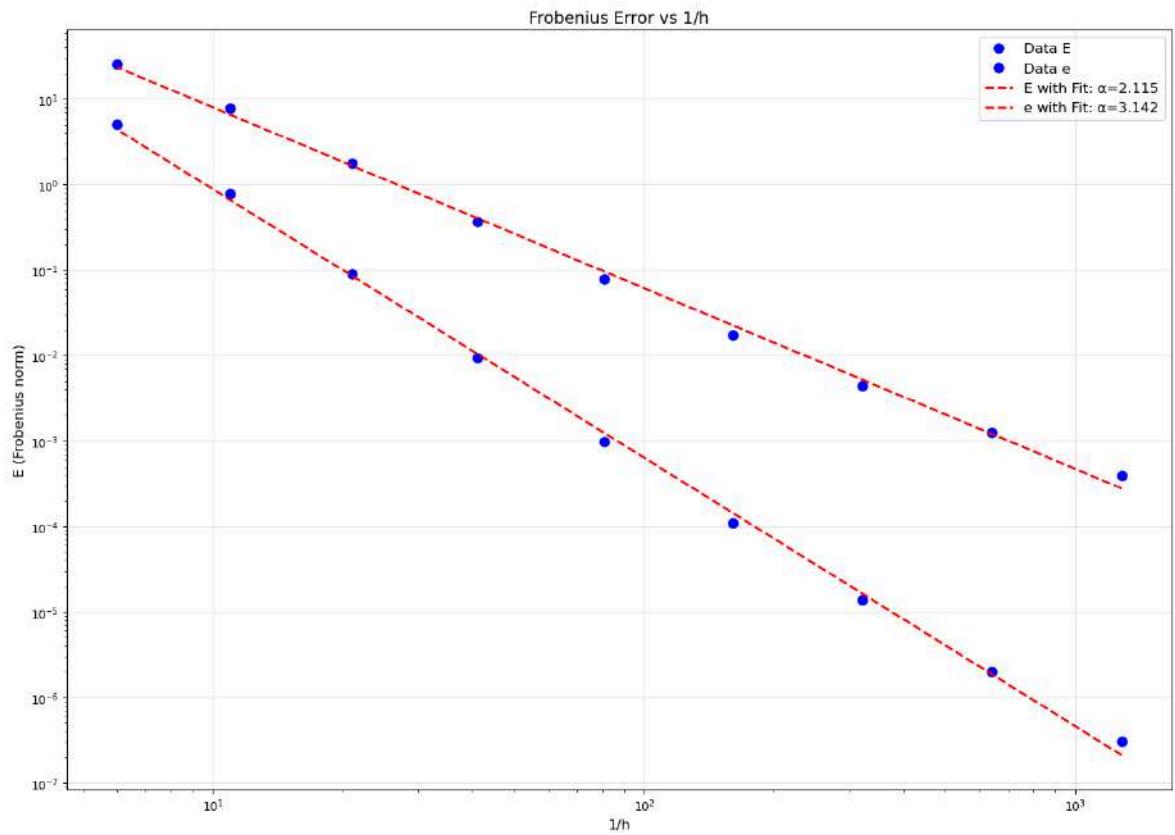
# Print convergence rates
print("The document was converted with Code-Format: https://code-format.com/
```

```

print("CONVERGENCE RATE ANALYSIS:")
print("=" * 80)
print(f"Frobenius norm (E): α = {α_E:.4f}")
print(f"Scaled norm (e): α = {α_e:.4f}")
return results, E_fit, e_fit, α_E, α_e
results_2, E_fit_2, e_fit_2, α_E2, α_e2=
fit_and_plot_convergence(fd_order=2)

```

N	h	E (Frobenius)	e (scaled)
Maximum error compared to N=1000 reference: 1.973921e+01			
5	1.666667e-01	2.553305e+01	5.106610e+00
Maximum error compared to N=1000 reference: 4.196089e+00			
10	9.090909e-02	7.849262e+00	7.849262e-01
Maximum error compared to N=1000 reference: 6.533160e-01			
20	4.761905e-02	1.795448e+00	8.977241e-02
Maximum error compared to N=1000 reference: 8.974813e-02			
40	2.439024e-02	3.726019e-01	9.315047e-03
Maximum error compared to N=1000 reference: 1.301626e-02			
80	1.234568e-02	7.762972e-02	9.703715e-04
Maximum error compared to N=1000 reference: 2.149269e-03			
160	6.211180e-03	1.740160e-02	1.087600e-04
Maximum error compared to N=1000 reference: 3.974418e-04			
320	3.115265e-03	4.397576e-03	1.374242e-05
Maximum error compared to N=1000 reference: 8.179131e-05			
640	1.560062e-03	1.257535e-03	1.964899e-06
Maximum error compared to N=1000 reference: 1.819485e-05			
1280	7.806401e-04	3.913961e-04	3.057782e-07
=====			
=====			
E (Frobenius) fit: E = 1043.646475 * h^2.1152			
e (scaled) fit: e = 1227.547608 * h^3.1425			



CONVERGENCE RATE ANALYSIS:

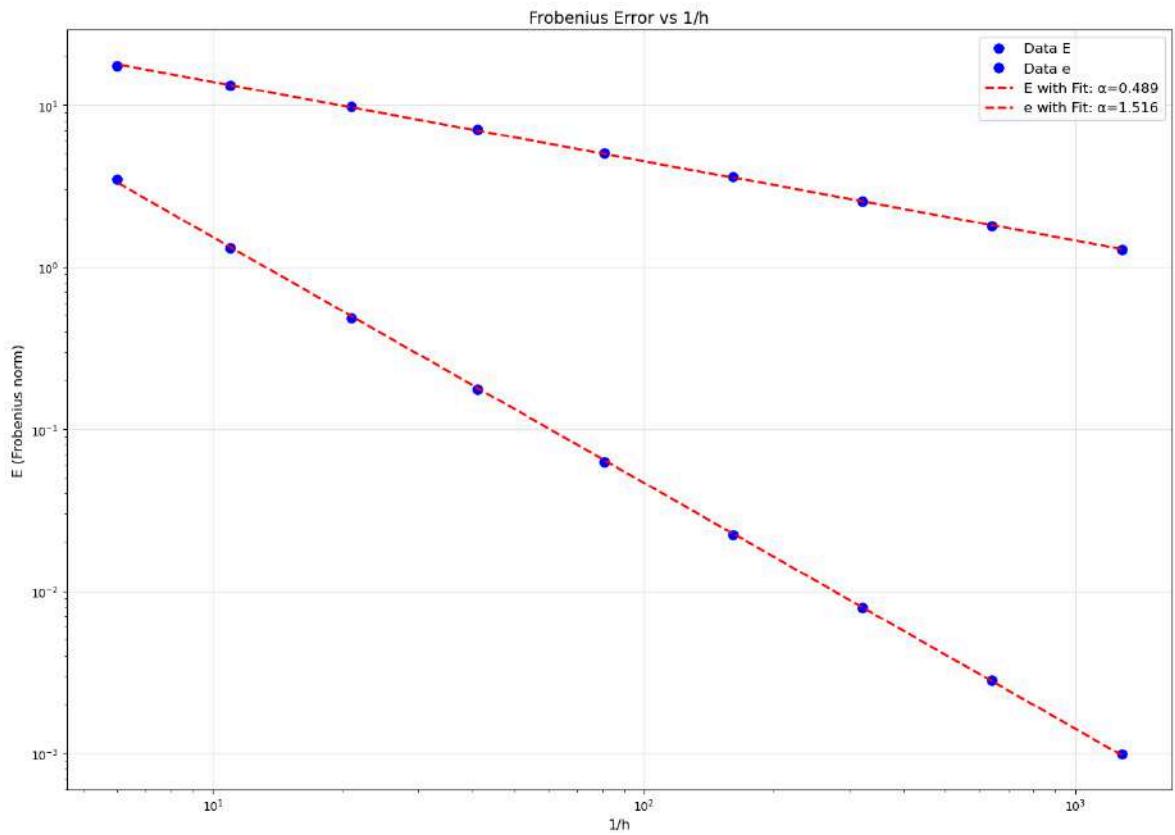
Frobenius norm (E): $\alpha = 2.1152$)

Scaled norm (e): $\alpha = 3.1425$

In [14]:

```
results_1, E_fit_1, e_fit_1, α_E1, α_e1 =
fit_and_plot_convergence(fd_order=1)
```

N	h	E (Frobenius)	e (scaled)
Maximum error compared to N=10000 reference: 1.315947e+01			
5	1.666667e-01	1.740636e+01	3.481271e+00
Maximum error compared to N=10000 reference: 7.177894e+00			
10	9.090909e-02	1.316567e+01	1.316567e+00
Maximum error compared to N=10000 reference: 3.759849e+00			
20	4.761905e-02	9.717364e+00	4.858682e-01
Maximum error compared to N=10000 reference: 1.925776e+00			
40	2.439024e-02	7.033490e+00	1.758373e-01
Maximum error compared to N=10000 reference: 9.747757e-01			
80	1.234568e-02	5.033981e+00	6.292476e-02
Maximum error compared to N=10000 reference: 4.904151e-01			
160	6.211180e-03	3.581529e+00	2.238456e-02
Maximum error compared to N=10000 reference: 2.459714e-01			
320	3.115265e-03	2.540388e+00	7.938712e-03
Maximum error compared to N=10000 reference: 1.231776e-01			
640	1.560062e-03	1.799122e+00	2.811129e-03
Maximum error compared to N=10000 reference: 6.163687e-02			
1280	7.806401e-04	1.273162e+00	9.946577e-04
=====			
=====			
E (Frobenius) fit: E = 42.687632 * h^0.4891			
e (scaled) fit: e = 50.209627 * h^1.5164			



CONVERGENCE RATE ANALYSIS:

Frobenius norm (E): $\alpha = 0.4891$)

Scaled norm (e): $\alpha = 1.5164$

In [15]:

```
fig, ax = plt.subplots(figsize=(10, 7))

# Plot order 2 results
ax.loglog(1/results_2['h'], results_2['E'], 'bo', linewidth=2,
           markersize=8,
           label='Order 2: E data')
ax.loglog(1/results_2['h'], results_2['e'], 'bs', linewidth=2,
           markersize=8,
           label='Order 2: e data', alpha=0.7)

# Plot order 1 results
ax.loglog(1/results_1['h'], results_1['E'], 'ro', linewidth=2,
           markersize=8,
           label='Order 1: E data')
ax.loglog(1/results_1['h'], results_1['e'], 'rs', linewidth=2,
           markersize=8,
           label='Order 1: e data', alpha=0.7)

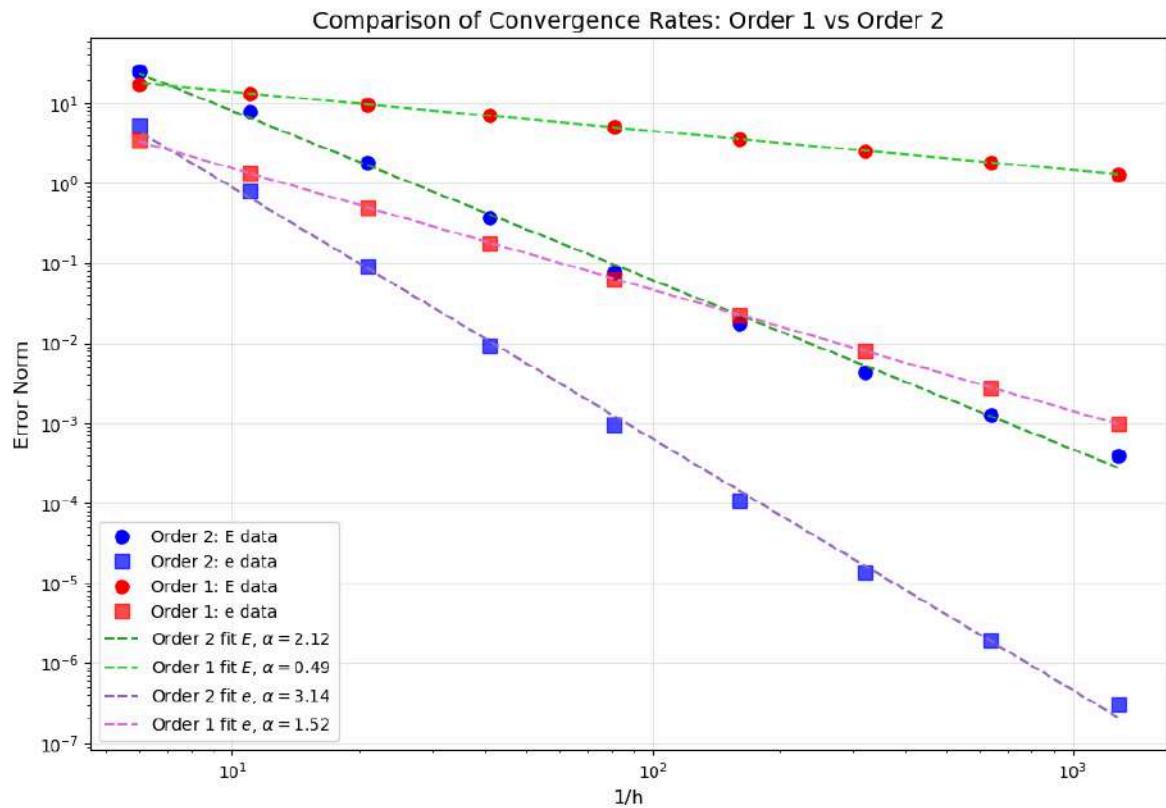
h_fit = np.logspace(np.log10(min(results_1['h'])),
                    np.log10(max(results_1['h'])), 100)
ax.loglog(1/h_fit, E_fit_2, color='tab:green', linestyle='--',
           linewidth=1.5,
           label=rf'Order 2 fit $E$, $\alpha=\{\alpha_E2:.2f\}$')

ax.loglog(1/h_fit, E_fit_1, color='limegreen', linestyle='--',
           linewidth=1.5,
           label=rf'Order 1 fit $E$, $\alpha=\{\alpha_E1:.2f\}$')

ax.loglog(1/h_fit, e_fit_2, color='tab:purple', linestyle='--',
           linewidth=1.5,
           label=rf'Order 2 fit $e$, $\alpha=\{\alpha_e2:.2f\}$')

ax.loglog(1/h_fit, e_fit_1, color='orchid', linestyle='--',
           linewidth=1.5,
           label=rf'Order 1 fit $e$, $\alpha=\{\alpha_e1:.2f\}$')

ax.set_xlabel('1/h', fontsize=12)
ax.set_ylabel('Error Norm', fontsize=12)
ax.set_title('Comparison of Convergence Rates: Order 1 vs Order 2',
             fontsize=14)
ax.grid(True, alpha=0.3)
ax.legend(fontsize=10)
plt.tight_layout()
plt.show()
```



Check with smaller step sizes:

In [16]:

```
results_2, E_fit_2, e_fit_2, alpha_E2, alpha_e2=
fit_and_plot_convergence(fd_order=2, N_values = [5*2**i for i in
range(5,10)])
results_1, E_fit_1, e_fit_1, alpha_E1, alpha_e1 =
fit_and_plot_convergence(fd_order=1, N_values = [5*2**i for i in
range(5,10)])
fig, ax = plt.subplots(figsize=(10, 7))

# Plot order 2 results
ax.loglog(1/results_2['h'], results_2['E'], 'bo', linewidth=2,
markersize=8,
           label='Order 2: E data')
ax.loglog(1/results_2['h'], results_2['e'], 'bs', linewidth=2,
markersize=8,
           label='Order 2: e data', alpha=0.7)

# Plot order 1 results
ax.loglog(1/results_1['h'], results_1['E'], 'ro', linewidth=2,
markersize=8,
           label='Order 1: E data')
ax.loglog(1/results_1['h'], results_1['e'], 'rs', linewidth=2,
markersize=8,
           label='Order 1: e data', alpha=0.7)

h_fit = np.logspace(np.log10(min(results_1['h'])),
np.log10(max(results_1['h'])), 100)
ax.loglog(1/h_fit, E_fit_2, color='tab:green', linestyle='--',
linewidth=1.5,
           label=rf'Order 2 fit $E$, $\alpha={\alpha_E2:.2f}$')

ax.loglog(1/h_fit, E_fit_1, color='limegreen', linestyle='--',
linewidth=1.5,
           label=rf'Order 1 fit $E$, $\alpha={\alpha_E1:.2f}$')

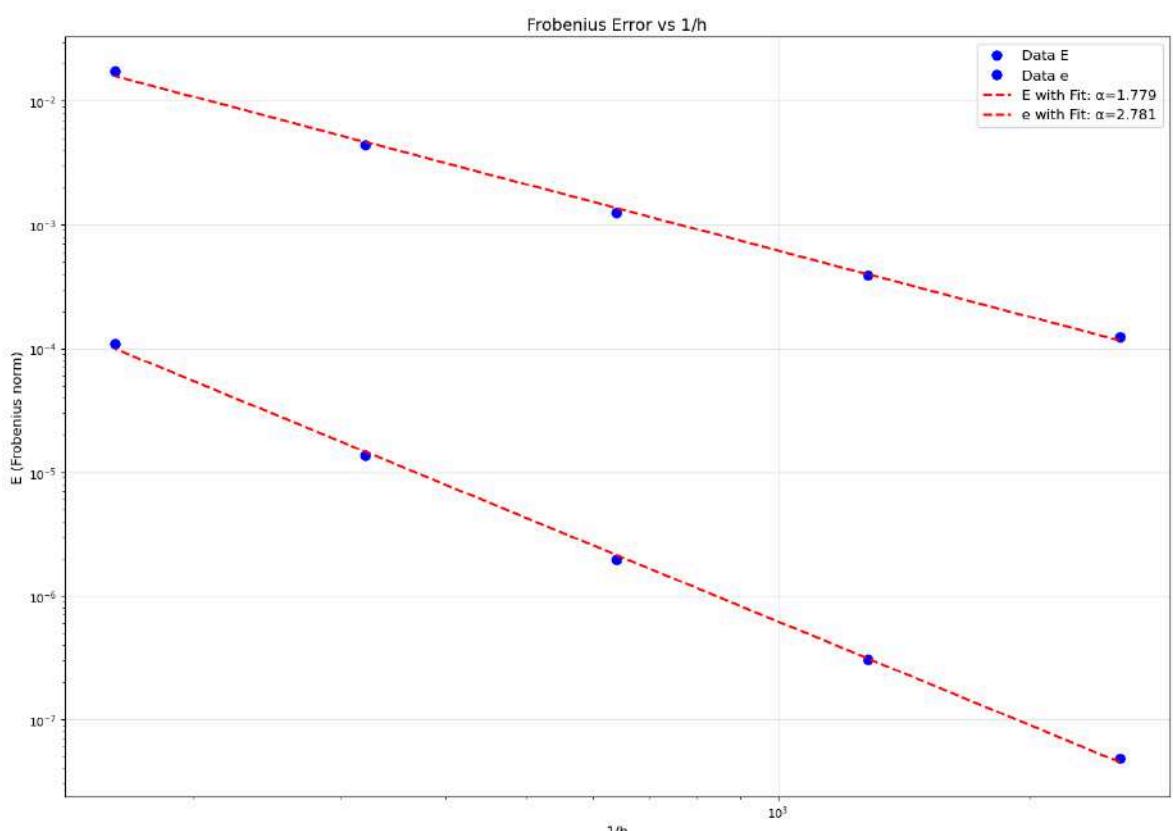
ax.loglog(1/h_fit, e_fit_2, color='tab:purple', linestyle='--',
linewidth=1.5,
           label=rf'Order 2 fit $e$, $\alpha={\alpha_e2:.2f}$')

ax.loglog(1/h_fit, e_fit_1, color='orchid', linestyle='--',
linewidth=1.5,
           label=rf'Order 1 fit $e$, $\alpha={\alpha_e1:.2f}$')

ax.set_xlabel('1/h', fontsize=12)
ax.set_ylabel('Error Norm', fontsize=12)
ax.set_title('Comparison of Convergence Rates: Order 1 vs Order 2',
fontsize=14)
ax.grid(True, alpha=0.5)
```

```
ax.legend(fontsize=10)
plt.tight_layout()
plt.show()
```

N	h	E (Frobenius)	e (scaled)
Maximum error compared to N=10000 reference: 2.149269e-03			
160	6.211180e-03	1.740160e-02	1.087600e-04
Maximum error compared to N=10000 reference: 3.974418e-04			
320	3.115265e-03	4.397576e-03	1.374242e-05
Maximum error compared to N=10000 reference: 8.179131e-05			
640	1.560062e-03	1.257535e-03	1.964899e-06
Maximum error compared to N=10000 reference: 1.819485e-05			
1280	7.806401e-04	3.913961e-04	3.057782e-07
Maximum error compared to N=10000 reference: 4.219866e-06			
2560	3.904725e-04	1.239707e-04	4.842606e-08
=====			
=====			
E (Frobenius) fit: E = 134.389496 * h^1.7792			
e (scaled) fit: e = 136.490499 * h^2.7812			



```
=====
=====
CONVERGENCE RATE ANALYSIS:
=====
=====
Frobenius norm (E): α = 1.7792)
Scaled norm (e):      α = 2.7812
      N          h          E (Frobenius)          e (scaled)

Maximum error compared to N=10000 reference: 4.904151e-01
    160 6.211180e-03      3.581529e+00      2.238456e-02

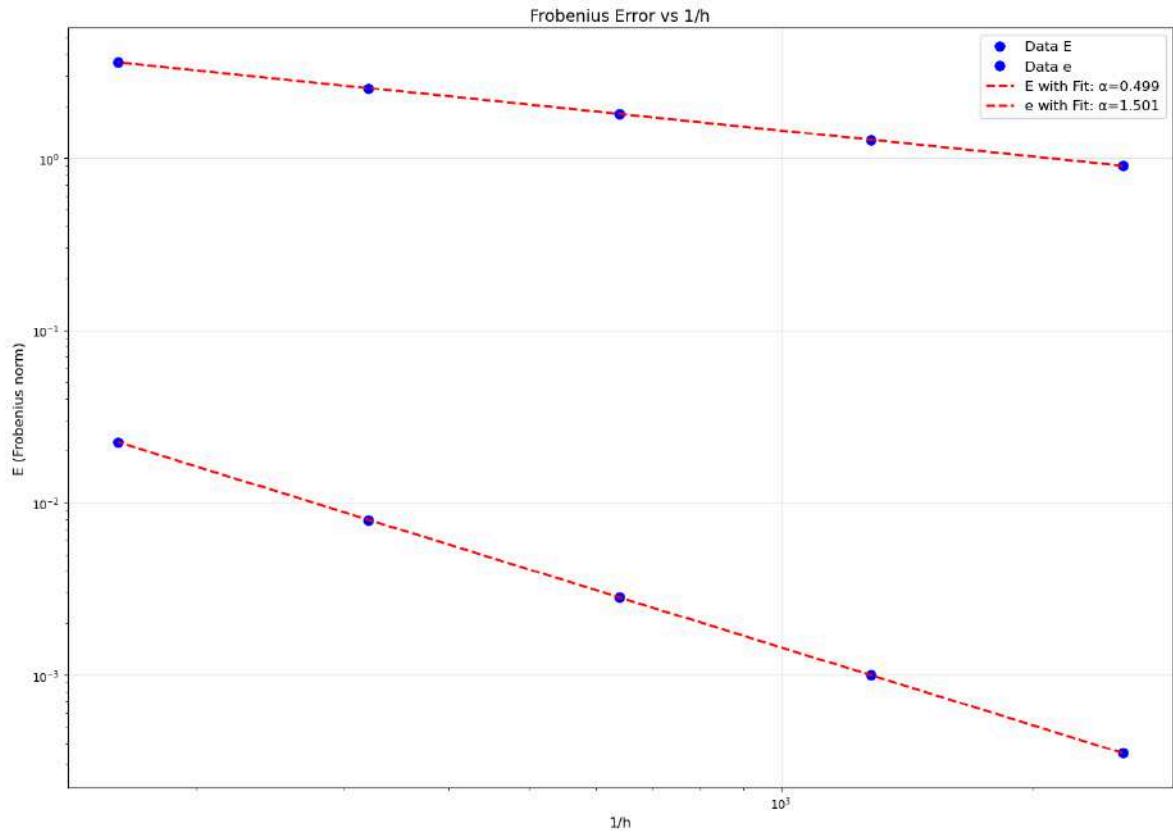
Maximum error compared to N=10000 reference: 2.459714e-01
    320 3.115265e-03      2.540388e+00      7.938712e-03

Maximum error compared to N=10000 reference: 1.231776e-01
    640 1.560062e-03      1.799122e+00      2.811129e-03

Maximum error compared to N=10000 reference: 6.163687e-02
    1280 7.806401e-04      1.273162e+00      9.946577e-04

Maximum error compared to N=10000 reference: 3.083046e-02
    2560 3.904725e-04      9.006098e-01      3.518007e-04
=====
=====

E (Frobenius) fit: E = 45.235386 * h^0.4990
e (scaled) fit:     e = 45.942581 * h^1.5010
```



=====

=====

=====

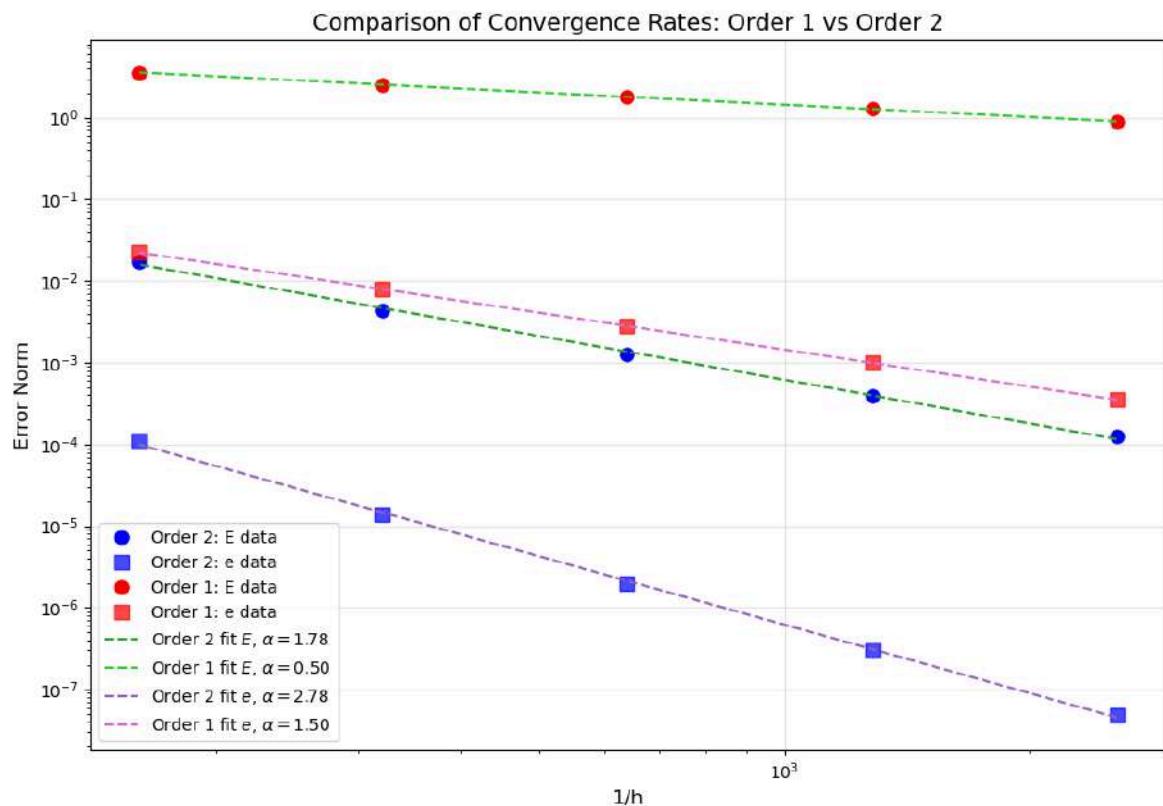
CONVERGENCE RATE ANALYSIS:

=====

=====

Frobenius norm (E): $\alpha = 0.4990$

Scaled norm (e): $\alpha = 1.5010$



Using a second-order one-sided finite difference approximation for the derivative at $x = 0$ the fitted convergence rates were

$$\alpha_E \approx 2.12, \quad \alpha_e \approx 3.14.$$

At a smaller step size,

$$\alpha_E \approx 1.78, \quad \alpha_e \approx 2.78.$$

The observed convergence rate for E is consistent with second-order accuracy, indicating that the overall discretization error is dominated by the interior second-order finite difference scheme together with the second-order boundary treatment. The error e converges at a significantly higher rate because it includes an additional factor proportional to the mesh size h . Specifically, since

$$e = \frac{E}{N}$$

and

$$N \sim \frac{1}{h},$$

it follows that

$$e \sim hE.$$

Consequently, e is expected to converge one order faster than E , which is consistent with the observed convergence rate of approximately 2.5. Note that these convergence rates are asymptotic, so at large step sizes the converge rates may not be exact; asymptotically, E will converge at 1.5 and e will converge at order 2.5. For more detail (on why 1.5 and 2.5 orders), please view problem 3.

Next, the second-order one-sided derivative at $x = 0$ was replaced with a first-order forward difference approximation. In this case, the fitted convergence rates were

$$\alpha_E \approx 0.49, \quad \alpha_e \approx 1.52.$$

We assume the pointwise discretization error satisfies

$$|u_h(x_i) - u(x_i)| = O(h^1).$$

per the first order discretization

Then the discrete ℓ^2 error is

$$E = \left(\sum_{i=1}^N O(h^2) \right)^{1/2} = O(\sqrt{N}, h^1) = O(h^{1/2}),$$

since

$$N \sim \frac{1}{h}.$$

Next, define the scaled error

$$e = \frac{1}{N} \|u_h - u\|_{\ell^2}.$$

Because

$$\frac{1}{N} \sim h,$$

we obtain

$$e = h, E.$$

Therefore,

$$e = O! \left(h^{1/2 + 1} \right) = O! \left(h^{3/2} \right).$$